

```
In [1]: import pandas as pd
from ast import literal_eval
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np
```

```
In [2]: import nltk
import re
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.model_selection import train_test_split
from nltk.stem import PorterStemmer
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\rezar\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\rezar\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\rezar\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

Load Data

```
In [68]: path = 'dataset/' #change to your data location
# Load downloaded data
df_convos = pd.read_csv(path+'/conversations.csv')
df_speakers = pd.read_csv(path+'/speakers.csv')
df_utts = pd.read_csv(path+'/utterances.csv')
df_cases = pd.read_json(path_or_buf=path+'/cases.jsonl', lines=True)
df_cases = df_cases[(df_cases['year'] >= 2011) & (df_cases['year'] <= 2018) & (df_cases['win_side'].isin([0,1]))]
```

```
In [69]: # count number win/lose cases
df_cases['win_side'].value_counts()
```

```
Out[69]: 1.0    400
0.0     201
Name: win_side, dtype: int64
```

```
In [70]: # combine text from all utterances in a conversation back into one string based on the conversation_id, coount how many utterances
utt_per_conv = df_utts.groupby('conversation_id')['text'].apply(lambda x: ' '.join(x)).reset_index()
utt_per_conv['num_utterances'] = df_utts.groupby('conversation_id')['text'].count().reset_index()['text']

# add the combined text to the conversations dataframe, merge on conversation_id in utt_per_conv and id in df_convo
df_convos_utt = df_convos.merge(utt_per_conv, left_on='id', right_on='conversation_id', how='left')
```

```
In [71]: # combine text from all conversation in a cases into one string based on the meta.case_id
conv_per_case = df_convos_utt.groupby('meta.case_id')['text'].apply(lambda x: ' '.join(x)).reset_index()
conv_per_case['num_conversations'] = df_convos_utt.groupby('meta.case_id')['text'].count().reset_index()['text']
conv_per_case['num_utterances'] = df_convos_utt.groupby('meta.case_id')['num_utterances'].sum().reset_index()['num_utterances']

# add the combined text case dataframe, merge on meta.case_id and id
df_cases_convo = df_cases.merge(conv_per_case, left_on='id', right_on='meta.case_id', how='left')
df_cases_convo.head(1)
```

```
Out[71]:
```

| | id | year | citation | title | petitioner | respondent | docket_no | court | decided_date | url | ... | win_side_detail |
|---|--------------|------|----------|---|--------------------------------------|---|-----------|---------------|--------------|---|-----|-----------------|
| 0 | 2011_11-1179 | 2011 | 567 US - | American Tradition Partnership, Inc. v. Bullock | American Tradition Partnership, Inc. | Steve Bullock, Attorney General of Montana, et... | 11-1179 | Roberts Court | Jun 25, 2012 | https://www.oyez.org/cases/2011/11-1179 | ... | 3.0 |

1 rows × 25 columns

```
In [72]: df_cases_convo.dropna(subset=['text'], inplace=True)
df_cases_convo.shape[0]
```

Out[72]: 521

```
In [73]: # transform to pd.to_datetime
df_cases_convo.decided_date = pd.to_datetime(df_cases_convo.decided_date)
```

Data Preprocessing

```
In [74]: # Cleaning the text
def preprocess_text(text):
    text = text.lower() # Lowercase the text
    text = re.sub('[^a-z]+', ' ', text) # Remove special characters and numbers
    text = re.sub(r'\b\w{1,3}\b', '', text) # Remove words with length less than 3
    words = nltk.word_tokenize(text) # Tokenize the text
    stop_words = set(stopwords.words('english')) # Remove stopwords
    words = [word for word in words if word not in stop_words]
    #Lemmatizer = WordNetLemmatizer() # Lemmatize the words comment because slow
    #words = [Lemmatizer.Lemmatize(word) for word in words]
    stemmer = PorterStemmer() # Stem the words
    words = [stemmer.stem(word) for word in words]
    text = ' '.join(words) # Reconstruct the text

    return text
```


```
In [75]: df_cases_convo.loc[:, 'text'] = df_cases_convo.loc[:, 'text'].apply(preprocess_text) #apply preprocess
df_cases_convo.to_csv('df_cases_clean.csv', index=False)
```

```
In [76]: text = pd.read_csv('df_cases_clean.csv')
text.head(1)
```

```
Out[76]:
```

| | id | year | citation | title | petitioner | respondent | docket_no | court | decided_date | url | ... | win_side_detail | scdt |
|---|-------------|------|----------|--------------------------|----------------|---------------|-----------|---------------|--------------|--|-----|-----------------|------|
| 0 | 2011_11-182 | 2011 | 567 US - | Arizona v. United States | Arizona et al. | United States | 11-182 | Roberts Court | 2012-06-25 | https://www.oyez.org/cases/2011/11-182 | ... | 7.0 | 2 |

1 rows × 25 columns



```
In [105... df_cases_convo.columns
```

```
Out[105]: Index(['id', 'year', 'citation', 'title', 'petitioner', 'respondent',
               'docket_no', 'court', 'decided_date', 'url', 'transcripts',
               'adv_sides_inferred', 'known_respondent_adv', 'advocates', 'win_side',
               'win_side_detail', 'scdb_docket_id', 'votes', 'votes_detail',
               'is_eq_divided', 'votes_side', 'meta.case_id', 'text',
               'num_conversations', 'num_utterances', 'start_date', 'develop_time'],
              dtype='object')
```

```
In [77]: # preprocess develop time
df_cases_convo.loc[:, 'start_date'] = df_cases_convo['transcripts'].apply(lambda x : re.findall(r'[A-Z][a-z]+ \d{2}, \d{4}', x)[0])
df_cases_convo.start_date = pd.to_datetime(df_cases_convo.start_date)
df_cases_convo.loc[:, 'develop_time'] = df_cases_convo.loc[:, 'decided_date'] - df_cases_convo.loc[:, 'start_date']
```

```
In [78]: # get party of the judges
def check_party_pc(x):
    rep_judge = ['j_clarence_thomas', 'j_anthony_m_kennedy', 'j_antonin_scalia', 'j_john_g_roberts_jr', 'j_samuel_a_alito_jr',
                 'j_john_paul_stevens', 'j_david_h_souter', 'j_william_h_rehnquist', 'j_neil_gorsuch', 'j_brett_m_kavanaugh']
    dem_judge = ['j_ruth_bader_ginsburg', 'j_stephen_g_breyer', 'j_sonia_sotomayor', 'j_elena_kagan']

    rep_ct = 0

    for judge in x:
        if judge in rep_judge:
            rep_ct += 1

    return rep_ct/len(x)
```

```
In [79]: # get rep_judge yes
def check_rep_j_y_pc(x):
    rep_judge = ['j_clarence_thomas', 'j_anthony_m_kennedy', 'j_antonin_scalia', 'j_john_g_roberts_jr', 'j_samuel_a_alito_jr',
                 'j_john_paul_stevens', 'j_david_h_souter', 'j_william_h_rehnquist', 'j_neil_gorsuch', 'j_brett_m_kavanaugh']
    dem_judge = ['j_ruth_bader_ginsburg', 'j_stephen_g_breyer', 'j_sonia_sotomayor', 'j_elena_kagan']
```

```

rep_y_ct = 0

for judge in x:
    if judge in rep_judge:
        if x[judge] > 0:
            rep_y_ct += 1

return rep_y_ct/len(x)

```

```

In [80]: # get dem_judge yes
def check_dem_j_y_pc(x):
    rep_judge = ['j__clarence_thomas', 'j__anthony_m_kennedy', 'j__antonin_scalia', 'j__john_g_roberts_jr', 'j__samuel_a_alito_jr',
                'j__john_paul_stevens', 'j__david_h_souter', 'j__william_h_rehnquist', 'j__neil_gorsuch', 'j__brett_m_kavanaugh']
    dem_judge = ['j__ruth_bader_ginsburg', 'j__stephen_g_breyer', 'j__sonia_sotomayor', 'j__elena_kagan']

    dem_y_ct = 0

    for judge in x:
        if judge in dem_judge:
            if x[judge] > 0:
                dem_y_ct += 1

    return dem_y_ct/len(x)

```

```

In [81]: def check_party(x):
    if x > 2009:
        return 0
    else:
        return 1

```

```

In [82]: # get M-F percentage in judges
def check_FM_jpc(x):
    male_judge = ['j__clarence_thomas',
                  'j__anthony_m_kennedy',
                  'j__antonin_scalia',
                  'j__john_g_roberts_jr',
                  'j__samuel_a_alito_jr',
                  'j__john_paul_stevens',
                  'j__david_h_souter',
                  'j__william_h_rehnquist',
                  'j__neil_gorsuch',
                  'j__brett_m_kavanaugh',
                  'j__stephen_g_breyer']
    female_judge = ['j__ruth_bader_ginsburg',
                    'j__sonia_sotomayor',
                    'j__elena_kagan']

    male_ct = 0
    for judge in x:
        if judge in male_judge:
            male_ct += 1

    return male_ct/len(x)

```

```

In [83]: # get rep_judge yes
def check_M_j_y_pc(x):
    male_judge = ['j__clarence_thomas',
                  'j__anthony_m_kennedy',
                  'j__antonin_scalia',
                  'j__john_g_roberts_jr',
                  'j__samuel_a_alito_jr',
                  'j__john_paul_stevens',
                  'j__david_h_souter',
                  'j__william_h_rehnquist',
                  'j__neil_gorsuch',
                  'j__brett_m_kavanaugh',
                  'j__stephen_g_breyer']
    female_judge = ['j__ruth_bader_ginsburg',
                    'j__sonia_sotomayor',
                    'j__elena_kagan']

    male_y_ct = 0
    for judge in x:
        if judge in male_judge:
            if x[judge] > 0:
                male_y_ct += 1

    return male_y_ct/len(x)

```

```

In [84]: # get rep_judge yes
def check_F_j_y_pc(x):
    male_judge = ['j__clarence_thomas',
                  'j__anthony_m_kennedy',

```

```

        'j__antonin_scalia',
        'j__john_g_roberts_jr',
        'j__samuel_a_alito_jr',
        'j__john_paul_stevens',
        'j__david_h_souter',
        'j__william_h_rehnquist',
        'j__neil_gorsuch',
        'j__brett_m_kavanaugh',
        'j__stephen_g_breyer']
female_judge = ['j__ruth_bader_ginsburg',
                'j__sonia_sotomayor',
                'j__elena_kagan']

female_y_ct = 0
for judge in x:
    if judge in female_judge:
        if x[judge] > 0:
            female_y_ct += 1
return female_y_ct/len(x)

```

```

In [85]: # get first name of speakers
df_test = df_cases_convo.loc[:, ['advocates']]

def get_side1_fstname(x):
    return list(x.keys())[0].split()[0]

def get_side0_fstname(x):
    try:
        return list(x.keys())[-1].split()[0]
    except:
        return list(x.keys())[-2].split()[0]

df_test.loc[:, 'side1_fstname'] = df_test['advocates'].apply(get_side1_fstname)
df_test.loc[:, 'side0_fstname'] = df_test['advocates'].apply(get_side0_fstname)

# read gender dataset
df_gender = pd.read_csv('dataset/name_gender_dataset.csv')
idx = df_gender.groupby(['Name'])['Probability'].idxmax()
df_gender = df_gender.loc[idx]

# join the gender dataset to predict gender of speakers
df_test = pd.merge(df_test, df_gender, how='left', left_on = 'side1_fstname', right_on = 'Name')
df_test = pd.merge(df_test, df_gender, how='left', left_on = 'side0_fstname', right_on = 'Name')

```

```

In [113... # only numbers can apply to Random Forest Model
df_rf = pd.DataFrame()
df_rf.loc[:, 'text_len'] = df_cases_convo['text'].apply(lambda x : len(x))
df_rf.loc[:, 'num_utterances'] = df_cases_convo['num_utterances']
df_rf.loc[:, 'win_side'] = df_cases_convo['win_side']
df_rf.loc[:, 'develop_time'] = df_cases_convo['develop_time'].apply(lambda x : x.days)
df_rf.loc[:, 'rep_jpc'] = df_cases_convo['votes_side'].apply(check_party_pc)
df_rf.loc[:, 'dem_jpc'] = 1 - df_rf.loc[:, 'rep_jpc']
df_rf.loc[:, 'rep_j_y_pc'] = df_cases_convo['votes_side'].apply(check_rep_j_y_pc)
df_rf.loc[:, 'dem_j_y_pc'] = df_cases_convo['votes_side'].apply(check_dem_j_y_pc)
df_rf.loc[:, 'party'] = df_cases_convo['year'].apply(check_party) # 1: rep, 0: dem
df_rf.loc[:, 'male_jpc'] = df_cases_convo['votes_side'].apply(check_FM_jpc)
df_rf.loc[:, 'female_jpc'] = 1 - df_rf.loc[:, 'male_jpc']
df_rf.loc[:, 'male_y_jpc'] = df_cases_convo['votes_side'].apply(check_M_j_y_pc)
df_rf.loc[:, 'female_y_jpc'] = df_cases_convo['votes_side'].apply(check_F_j_y_pc)
df_rf.loc[:, 'text'] = df_cases_convo['text']
df_rf.loc[:, 'num_conversations'] = df_cases_convo['num_conversations']

# reset the index
df_rf = df_rf.reset_index(drop=True)
df_rf.loc[:, 'side1_gender'] = df_test['Gender_x'].apply(lambda x: 0 if(x == 'F') else 1)
df_rf.loc[:, 'side0_gender'] = df_test['Gender_y'].apply(lambda x: 0 if(x == 'F') else 1)

```

```

In [88]: df_rf.to_csv('dataset/df_rf.csv', index=False)

```

```

In [3]: df_rf = pd.read_csv('dataset/df_rf.csv')

```

Additional Graphs for Checkpoint 1

```

In [4]: # Histogram of text length, number of utterances, and development time, and number of conversations for win_side = 1
fig, ax = plt.subplots(2, 2, figsize=(10, 10))
ax[0, 0].hist(df_rf[df_rf['win_side'] == 1]['text_len'], bins=50)
ax[0, 0].set_title('Histogram of Text Length')
ax[0, 1].hist(df_rf[df_rf['win_side'] == 1]['num_utterances'], bins=50)
ax[0, 1].set_title('Histogram of Number of Utterances')
ax[1, 0].hist(df_rf[df_rf['win_side'] == 1]['develop_time'], bins=50)

```

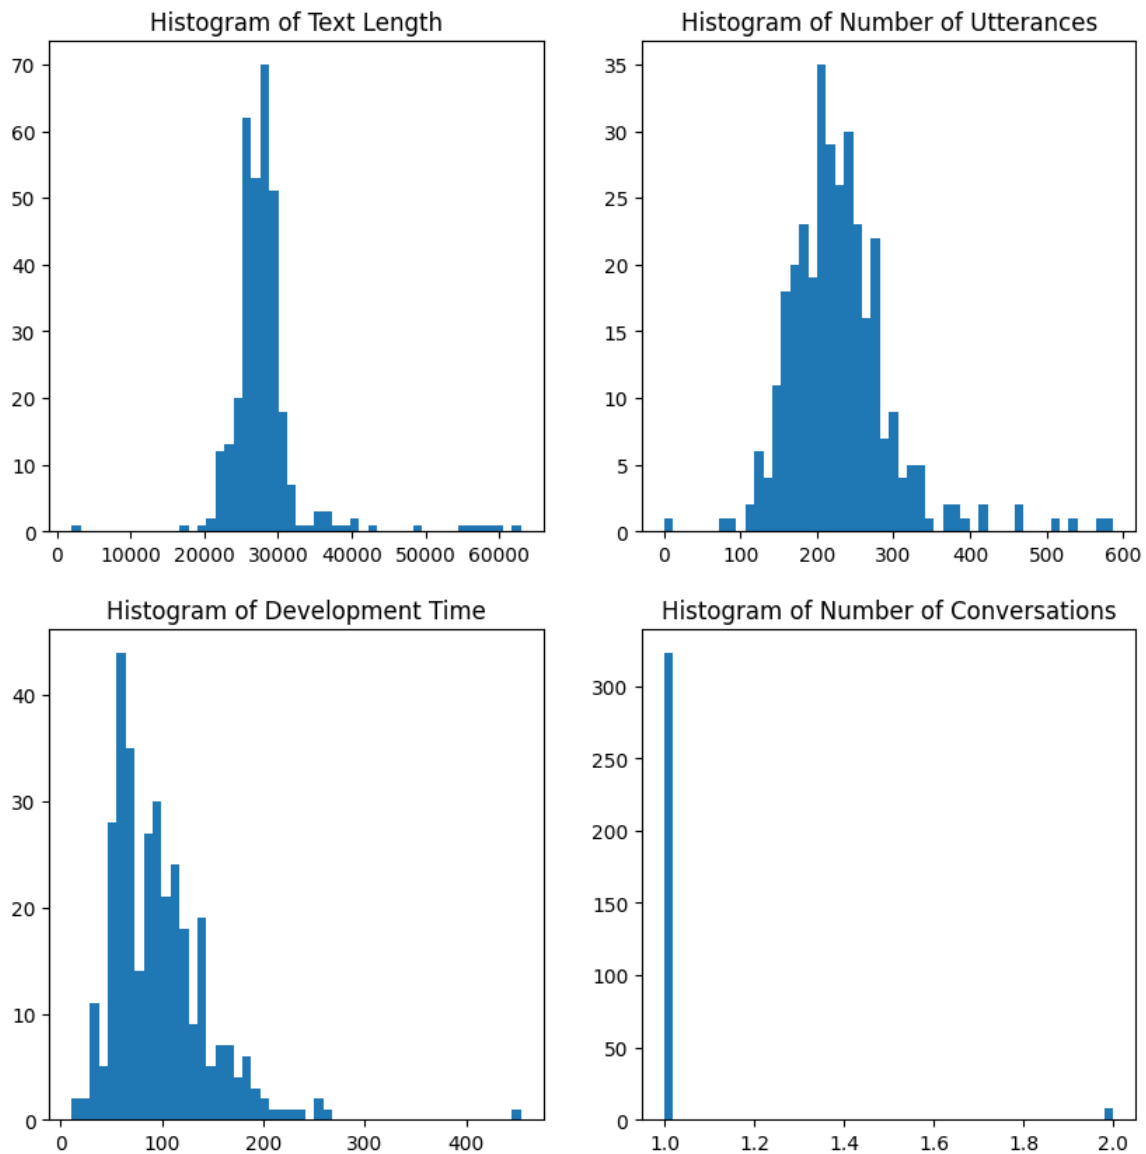
```

ax[1, 0].set_title('Histogram of Development Time')
ax[1, 1].hist(df_rf[df_rf['win_side'] == 1]['num_conversations'], bins=50)
ax[1, 1].set_title('Histogram of Number of Conversations')
fig.suptitle('Histograms for Petitioner Wins', fontsize=16)
plt.savefig('histogram1.png')

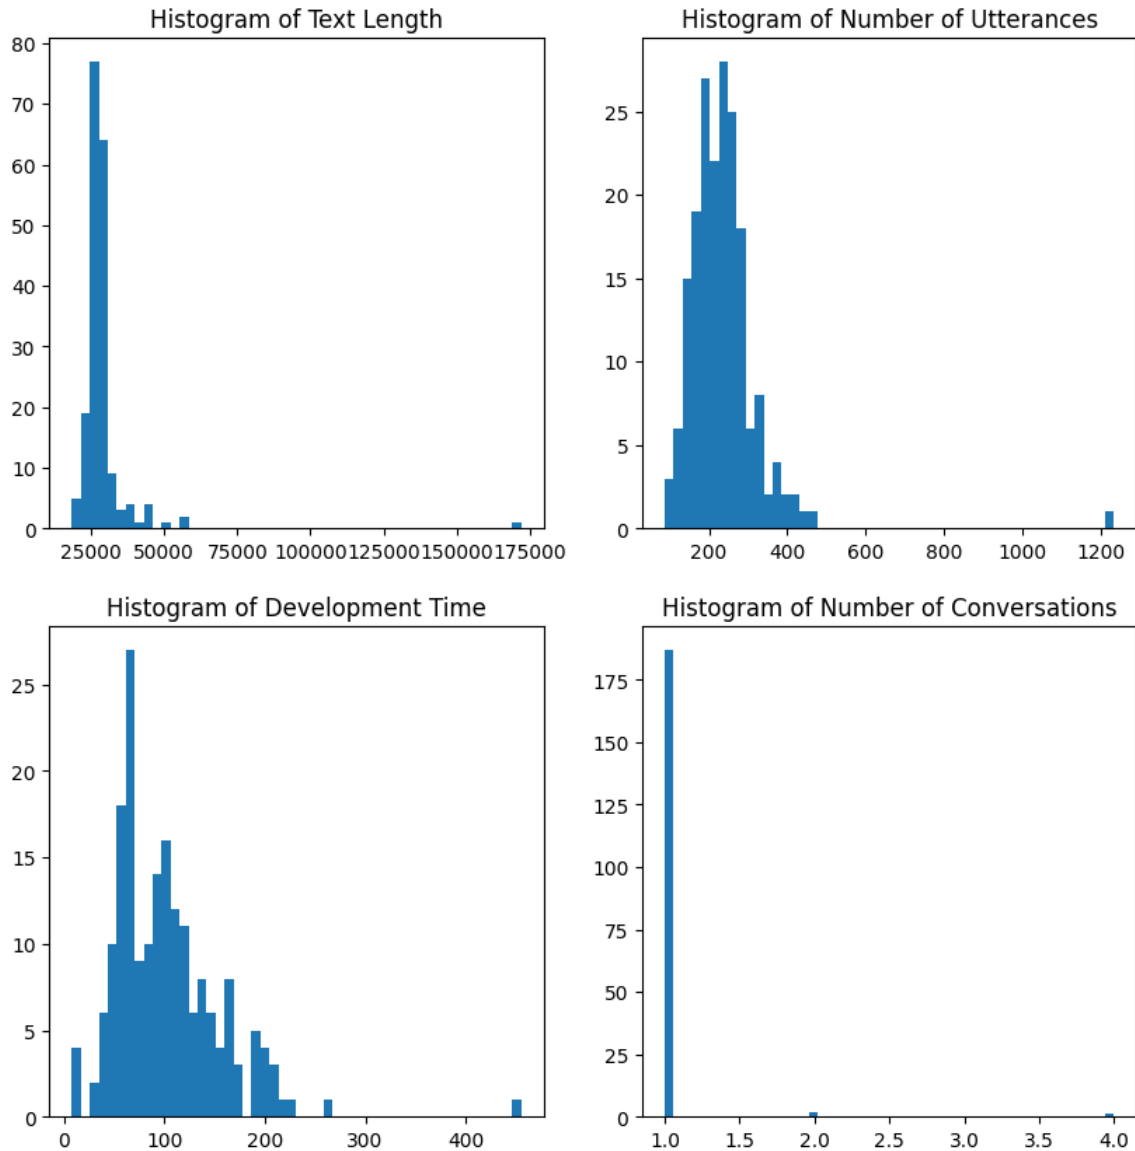
# Histogram of text length, number of utterances, and development time, and number of conversations for win_side = 0
fig, ax = plt.subplots(2, 2, figsize=(10, 10))
ax[0, 0].hist(df_rf[df_rf['win_side'] == 0]['text_len'], bins=50)
ax[0, 0].set_title('Histogram of Text Length')
ax[0, 1].hist(df_rf[df_rf['win_side'] == 0]['num_utterances'], bins=50)
ax[0, 1].set_title('Histogram of Number of Utterances')
ax[1, 0].hist(df_rf[df_rf['win_side'] == 0]['develop_time'], bins=50)
ax[1, 0].set_title('Histogram of Development Time')
ax[1, 1].hist(df_rf[df_rf['win_side'] == 0]['num_conversations'], bins=50)
ax[1, 1].set_title('Histogram of Number of Conversations')
fig.suptitle('Histograms for Respondent Wins', fontsize=16)
plt.savefig('histogram0.png')

```

Histograms for Petitioner Wins



Histograms for Respondent Wins



Baseline

```
In [5]: # Calculate The Baseline for Accuracy, Precision, Recall, F1
accuracy = df_rf['win_side'].value_counts()[1]/df_rf['win_side'].shape[0]
print('Accuracy: ', accuracy)
```

Accuracy: 0.6353166986564299

Model Selection and Vectorize

```
In [45]: from sklearn.linear_model import LogisticRegression, Perceptron, SGDClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, \
    f1_score, make_scorer, precision_score, recall_score, ConfusionMatrixDisplay, roc_auc_score, roc_curve, auc

def Classifier(X_train, X_test, y_train, y_test):

    # Train and evaluate the classifiers
    classifiers = {
        "Logistic Regression": LogisticRegression(max_iter=1000),
```

```

"Naive Bayes": MultinomialNB(),
"Linear SVC": LinearSVC(),
"Random Forest": RandomForestClassifier(),
"Perceptron": Perceptron(),
"KNN": KNeighborsClassifier(n_neighbors=7)

}

results = []

for classifier_name, classifier in classifiers.items():

    # Train the classifier
    classifier.fit(X_train, y_train)

    # Make predictions on the test set
    y_pred = classifier.predict(X_test)

    # Add the scores to the results dictionary
    results.append({
        'classifier': classifier_name,
        'accuracy': accuracy_score(y_test, y_pred),
        'f1': f1_score(y_test, y_pred),
        'precision': precision_score(y_test, y_pred),
        'recall': recall_score(y_test, y_pred),
        'True Negative Rate': confusion_matrix(y_test, y_pred)[0][0]/(confusion_matrix(y_test, y_pred)[0][0]+confusion_matrix(
        'True Positive Rate': confusion_matrix(y_test, y_pred)[1][1]/(confusion_matrix(y_test, y_pred)[1][1]+confusion_matrix(
        'auc': roc_auc_score(y_test, y_pred)
    })

    # Make a confusion matrix
    print(f"Confusion Matrix for {classifier_name}:")

    if classifier_name == "Logistic Regression":
        print("Coefficients: \n", classifier.coef_)

    cm = confusion_matrix(y_test, y_pred, labels=y_test.unique())
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=y_test.unique())
    fig, ax = plt.subplots(figsize=(3, 3))
    ax.set_title(f"CM {classifier_name}:")
    disp.plot(ax=ax)
    plt.show()
return pd.DataFrame(results)

```

```

In [5]: def Vectorize(vectorizer, X, y):
        X = vectorizer.fit_transform(X)
        y = y
        return X, y

```

Using 'Text' as the feature

Text as Features and using TF-IDF Vectorizer

```

In [10]: # Vectorize the text using TF-IDF
vectorizer = TfidfVectorizer(min_df=5, max_df=0.7)
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)

```


Confusion Matrix for Logistic Regression:

Coefficients:


```

[[ 0.11731986  0.04298986 -0.00358675 ... -0.0470672  0.00655914
  0.02283127]]


```




Confusion Matrix for Naive Bayes:



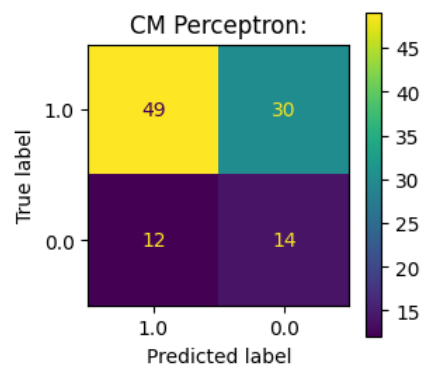
Confusion Matrix for Linear SVC:




Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:




Out[10]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.714286 | 0.825581 | 0.763441 | 0.898734 | 0.153846 | 0.898734 | 0.526290 |
| 1 | Naive Bayes | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 2 | Linear SVC | 0.571429 | 0.697987 | 0.742857 | 0.658228 | 0.307692 | 0.658228 | 0.482960 |
| 3 | Random Forest | 0.733333 | 0.840909 | 0.762887 | 0.936709 | 0.115385 | 0.936709 | 0.526047 |
| 4 | Perceptron | 0.600000 | 0.700000 | 0.803279 | 0.620253 | 0.538462 | 0.620253 | 0.579357 |
| 5 | KNN | 0.600000 | 0.723684 | 0.753425 | 0.696203 | 0.307692 | 0.696203 | 0.501947 |


Text as Features and using Count Vectorizer

```
In [9]: # Vectorize the text using CountVectorizer
vectorizer = CountVectorizer(min_df=5, max_df=0.8)
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```

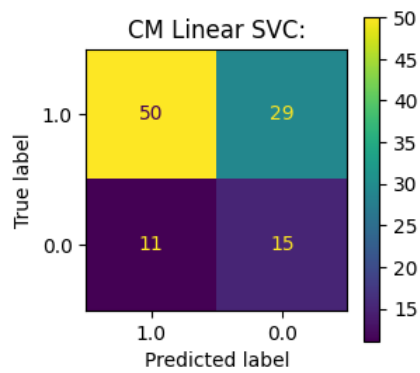
Confusion Matrix for Logistic Regression:




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:



Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Out[9]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.619048 | 0.726027 | 0.791045 | 0.670886 | 0.461538 | 0.670886 | 0.566212 |
| 1 | Naive Bayes | 0.561905 | 0.684932 | 0.746269 | 0.632911 | 0.346154 | 0.632911 | 0.489533 |
| 2 | Linear SVC | 0.619048 | 0.714286 | 0.819672 | 0.632911 | 0.576923 | 0.632911 | 0.604917 |
| 3 | Random Forest | 0.714286 | 0.831461 | 0.747475 | 0.936709 | 0.038462 | 0.936709 | 0.487585 |
| 4 | Perceptron | 0.638095 | 0.739726 | 0.805970 | 0.683544 | 0.500000 | 0.683544 | 0.591772 |

Text as Features and using TF-IDF Vectorizer and using Over and Under Sampling

```
In [11]: # USING IMBLEARN
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler

vectorizer = TfidfVectorizer(min_df=5, max_df=0.8)
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)

# Resample the training data
print('--OVERSAMPLING--')
ros = SMOTE(random_state=0)
X_train_resampled, y_train_resampled = ros.fit_resample(X_train, y_train)
classifier = Classifier(X_train_resampled, X_test, y_train_resampled, y_test)
print(classifier)

print('--UNDERSAMPLING--')
```

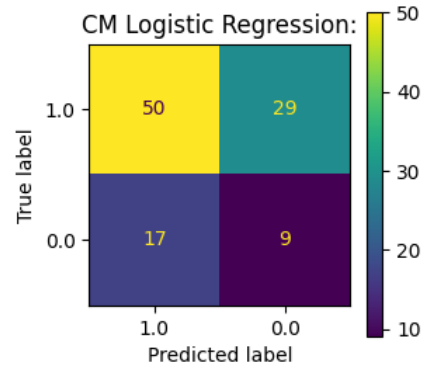
```

ros = RandomUnderSampler(random_state=0)
X_train_resampled, y_train_resampled = ros.fit_resample(X_train, y_train)
classifier = Classifier(X_train_resampled, X_test, y_train_resampled, y_test)
print(classifier)


```

--OVERSAMPLING--


Confusion Matrix for Logistic Regression:



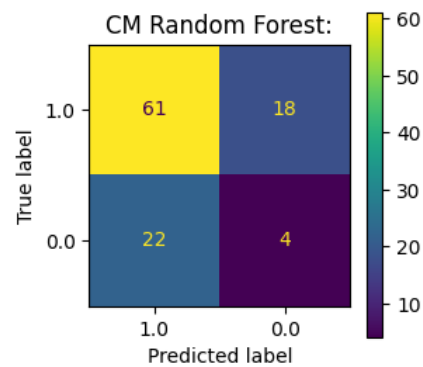
Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:



Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:




| | classifier | accuracy | f1 | precision | recall | \ |
|---|---------------------|----------|----------|-----------|----------|---|
| 0 | Logistic Regression | 0.561905 | 0.684932 | 0.746269 | 0.632911 | |
| 1 | Naive Bayes | 0.542857 | 0.661972 | 0.746032 | 0.594937 | |
| 2 | Linear SVC | 0.571429 | 0.680851 | 0.774194 | 0.607595 | |
| 3 | Random Forest | 0.619048 | 0.753086 | 0.734940 | 0.772152 | |
| 4 | Perceptron | 0.580952 | 0.681159 | 0.796610 | 0.594937 | |


| | True Negative Rate | True Positive Rate | auc |
|---|--------------------|--------------------|----------|
| 0 | 0.346154 | 0.632911 | 0.489533 |
| 1 | 0.384615 | 0.594937 | 0.489776 |
| 2 | 0.461538 | 0.607595 | 0.534567 |
| 3 | 0.153846 | 0.772152 | 0.462999 |
| 4 | 0.538462 | 0.594937 | 0.566699 |

--UNDERSAMPLING--

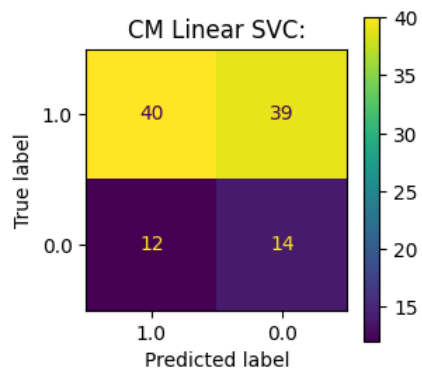
Confusion Matrix for Logistic Regression:




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:



Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



| | classifier | accuracy | f1 | precision | recall | \ |
|---|---------------------|----------|----------|-----------|----------|---|
| 0 | Logistic Regression | 0.466667 | 0.575758 | 0.716981 | 0.481013 | |
| 1 | Naive Bayes | 0.476190 | 0.592593 | 0.714286 | 0.506329 | |
| 2 | Linear SVC | 0.514286 | 0.610687 | 0.769231 | 0.506329 | |
| 3 | Random Forest | 0.523810 | 0.621212 | 0.773585 | 0.518987 | |
| 4 | Perceptron | 0.495238 | 0.589147 | 0.760000 | 0.481013 | |

| | True Negative Rate | True Positive Rate | auc |
|---|--------------------|--------------------|----------|
| 0 | 0.423077 | 0.481013 | 0.452045 |
| 1 | 0.384615 | 0.506329 | 0.445472 |
| 2 | 0.538462 | 0.506329 | 0.522395 |
| 3 | 0.538462 | 0.518987 | 0.528724 |
| 4 | 0.538462 | 0.481013 | 0.509737 |

Try using k-fold cross validation

```
In [33]: from sklearn.model_selection import cross_val_score

def Classifier_kfold(X, y):

    classifiers = {
        "Logistic Regression": LogisticRegression(max_iter=1000),
        "Naive Bayes": MultinomialNB(),
        "Linear SVC": LinearSVC(),
        "Random Forest": RandomForestClassifier(),
        "Perceptron": Perceptron(),
    }

    results = []
```

```

for classifier_name, classifier in classifiers.items():

    # Perform k-fold cross-validation
    scores = cross_val_score(classifier, X, y, cv=5, scoring='accuracy') # 5-fold cross-validation
    # Add the scores to the results list
    results.append({
        'classifier': classifier_name,
        'mean_f1': scores.mean(),
        'std_f1': scores.std(),
    })

    print(f"f1 for {classifier_name}: {scores.mean()} (+/- {scores.std() * 2})")

return pd.DataFrame(results)

# Vectorize the text using TF-IDF
vectorizer = TfidfVectorizer(min_df=5, max_df=0.7)
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
Classifier_kfold(X, y)

```

```

f1 for Logistic Regression: 0.6237912087912088 (+/- 0.008850077622725633)
f1 for Naive Bayes: 0.6353113553113553 (+/- 0.002783882783882774)
f1 for Linear SVC: 0.6083333333333334 (+/- 0.06005147912118134)
f1 for Random Forest: 0.6276373626373626 (+/- 0.03304313560378509)
f1 for Perceptron: 0.5853663003663004 (+/- 0.054826294118804415)

```

Out[33]:

| | classifier | mean_f1 | std_f1 |
|---|---------------------|----------|----------|
| 0 | Logistic Regression | 0.623791 | 0.004425 |
| 1 | Naive Bayes | 0.635311 | 0.001392 |
| 2 | Linear SVC | 0.608333 | 0.030026 |
| 3 | Random Forest | 0.627637 | 0.016522 |
| 4 | Perceptron | 0.585366 | 0.027413 |

Using ngrams

```

In [11]: # Vectorize the text using TF-IDF and bigrams
vectorizer = TfidfVectorizer(min_df=5, max_df=0.7, ngram_range=(2, 2))
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)

```

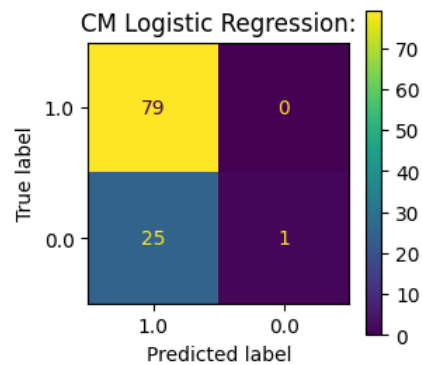
Confusion Matrix for Logistic Regression:

Coefficients:


```

[[ 0.01914151  0.05325384 -0.01408864 ... -0.01701203 -0.00579603
 -0.16882371]]


```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:



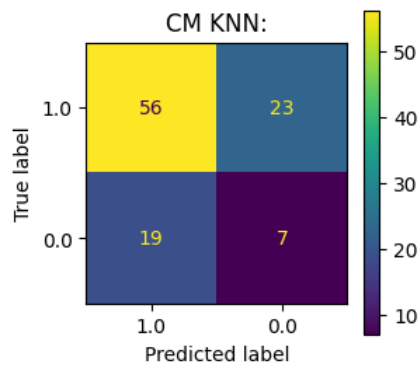
Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:



Out[11]:

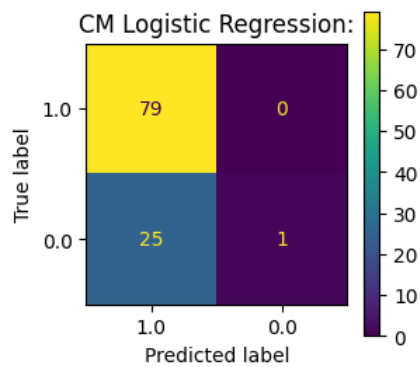
| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.647619 | 0.770186 | 0.756098 | 0.784810 | 0.230769 | 0.784810 | 0.507790 |
| 3 | Random Forest | 0.723810 | 0.837989 | 0.750000 | 0.949367 | 0.038462 | 0.949367 | 0.493914 |
| 4 | Perceptron | 0.600000 | 0.716216 | 0.768116 | 0.670886 | 0.384615 | 0.670886 | 0.527751 |
| 5 | KNN | 0.600000 | 0.727273 | 0.746667 | 0.708861 | 0.269231 | 0.708861 | 0.489046 |

```
In [12]: # Vectorize the text using TF-IDF: Trigrams
vectorizer = TfidfVectorizer(min_df=5, max_df=0.7, ngram_range=(3, 3))
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```

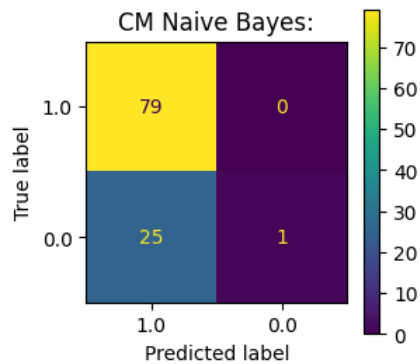
Confusion Matrix for Logistic Regression:

Coefficients:


```
[[-0.09237814  0.07358038  0.09318226 ...  0.07050802  0.00746334
 -0.24805484]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:




Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:



Out[12]:


| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 2 | Linear SVC | 0.628571 | 0.766467 | 0.727273 | 0.810127 | 0.076923 | 0.810127 | 0.443525 |
| 3 | Random Forest | 0.685714 | 0.811429 | 0.739583 | 0.898734 | 0.038462 | 0.898734 | 0.468598 |
| 4 | Perceptron | 0.561905 | 0.697368 | 0.726027 | 0.670886 | 0.230769 | 0.670886 | 0.450828 |
| 5 | KNN | 0.619048 | 0.746835 | 0.746835 | 0.746835 | 0.230769 | 0.746835 | 0.488802 |

```
In [13]: # Vectorize the text using TF-IDF: range 1 to 3 grams
vectorizer = TfidfVectorizer(min_df=5, max_df=0.7, ngram_range=(1, 3))
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```


Confusion Matrix for Logistic Regression:

Coefficients:


```
[[ 0.09887798  0.01125257  0.02364326 ...  0.02447635 -0.08093888
 -0.04026498]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:




Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:



| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.714286 | 0.827586 | 0.757895 | 0.911392 | 0.115385 | 0.911392 | 0.513389 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.571429 | 0.709677 | 0.723684 | 0.696203 | 0.192308 | 0.696203 | 0.444255 |
| 3 | Random Forest | 0.704762 | 0.826816 | 0.740000 | 0.936709 | 0.000000 | 0.936709 | 0.468354 |
| 4 | Perceptron | 0.600000 | 0.720000 | 0.760563 | 0.683544 | 0.346154 | 0.683544 | 0.514849 |
| 5 | KNN | 0.609524 | 0.732026 | 0.756757 | 0.708861 | 0.307692 | 0.708861 | 0.508277 |

```
In [70]: # Calculate The Baseline based on y_test
accuracy = y_test.value_counts()[1]/y_test.shape[0]
print('Accuracy: ', accuracy)
```

Accuracy: 0.7523809523809524

Using 'Text' + rep_pct + male_jpc + develop_time as the feature


```
In [14]: tfidf = TfidfVectorizer(min_df=5, max_df=0.8, ngram_range=(2, 2))
X_text = tfidf.fit_transform(df_rf['text'])
X_rep_pct = np.array(df_rf['rep_jpc']).reshape(-1, 1)
X_male_jpc = np.array(df_rf['male_jpc']).reshape(-1, 1)
X_dev_time = np.array(df_rf['develop_time']).reshape(-1, 1)
X_party = np.array(df_rf['party']).reshape(-1, 1)
y = df_rf['win_side']
# Horizontally stack the TF-IDF other features
X = np.hstack((X_text.toarray(), X_rep_pct, X_male_jpc, X_dev_time))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```


Confusion Matrix for Logistic Regression:

Coefficients:

```
[[ 0.01900358  0.05249416 -0.01398821 ...  0.01246409 -0.08228093
 -0.00045653]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\svm_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(




Confusion Matrix for Random Forest:




Confusion Matrix for Perceptron:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))



Confusion Matrix for KNN:



Out[14]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 3 | Random Forest | 0.733333 | 0.846154 | 0.747573 | 0.974684 | 0.000000 | 0.974684 | 0.487342 |
| 4 | Perceptron | 0.247619 | 0.000000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 | 0.500000 |
| 5 | KNN | 0.685714 | 0.800000 | 0.767442 | 0.835443 | 0.230769 | 0.835443 | 0.533106 |

Using 'Text' + rep_pct + male_jpc as the feature


```
In [15]: tfidf = TfidfVectorizer(min_df=5, max_df=0.8, ngram_range=(2, 2))
X_text = tfidf.fit_transform(df_rf['text'])
X_rep_pct = np.array(df_rf['rep_jpc']).reshape(-1, 1)
X_male_jpc = np.array(df_rf['male_jpc']).reshape(-1, 1)
X_dev_time = np.array(df_rf['develop_time']).reshape(-1, 1)
X_party = np.array(df_rf['party']).reshape(-1, 1)
y = df_rf['win_side']
# Horizontally stack the TF-IDF other features
X = np.hstack((X_text.toarray(), X_rep_pct, X_dev_time))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```


Confusion Matrix for Logistic Regression:

Coefficients:

```
[[ 0.01901738  0.05252966 -0.01400972 ... -0.16520139  0.01492379
 -0.00045276]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\svm_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(




Confusion Matrix for Random Forest:




Confusion Matrix for Perceptron:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))



Confusion Matrix for KNN:



Out[15]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 3 | Random Forest | 0.695238 | 0.818182 | 0.742268 | 0.911392 | 0.038462 | 0.911392 | 0.474927 |
| 4 | Perceptron | 0.247619 | 0.000000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 | 0.500000 |
| 5 | KNN | 0.685714 | 0.800000 | 0.767442 | 0.835443 | 0.230769 | 0.835443 | 0.533106 |

Using 'Text' + male_jpc + develop_time as the feature


```
In [16]: tfidf = TfidfVectorizer(min_df=5, max_df=0.8, ngram_range=(2, 2))
X_text = tfidf.fit_transform(df_rf['text'])
X_rep_pct = np.array(df_rf['rep_jpc']).reshape(-1, 1)
X_male_jpc = np.array(df_rf['male_jpc']).reshape(-1, 1)
X_dev_time = np.array(df_rf['develop_time']).reshape(-1, 1)
X_party = np.array(df_rf['party']).reshape(-1, 1)
y = df_rf['win_side']
# Horizontally stack the TF-IDF other features
X = np.hstack((X_text.toarray(), X_male_jpc, X_dev_time))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```


Confusion Matrix for Logistic Regression:

Coefficients:

```
[[ 0.01866398  0.05256018 -0.01407685 ... -0.16448145 -0.03192788
 -0.00044994]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\svm_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
 warnings.warn(
 C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
 _warn_prf(average, modifier, msg_start, len(result))




Confusion Matrix for Random Forest:




Confusion Matrix for Perceptron:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
 _warn_prf(average, modifier, msg_start, len(result))



Confusion Matrix for KNN:



Out[16]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.247619 | 0.000000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 | 0.500000 |
| 3 | Random Forest | 0.704762 | 0.826816 | 0.740000 | 0.936709 | 0.000000 | 0.936709 | 0.468354 |
| 4 | Perceptron | 0.247619 | 0.000000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 | 0.500000 |
| 5 | KNN | 0.685714 | 0.800000 | 0.767442 | 0.835443 | 0.230769 | 0.835443 | 0.533106 |

Using 'Text' + rep_jpc as the feature


```
In [17]: tfidf = TfidfVectorizer(min_df=5, max_df=0.8, ngram_range=(2, 2))
X_text = tfidf.fit_transform(df_rf['text'])
X_rep_pct = np.array(df_rf['rep_jpc']).reshape(-1, 1)
X_male_jpc = np.array(df_rf['male_jpc']).reshape(-1, 1)
X_dev_time = np.array(df_rf['develop_time']).reshape(-1, 1)
X_party = np.array(df_rf['party']).reshape(-1, 1)
y = df_rf['win_side']
# Horizontally stack the TF-IDF other features
X = np.hstack((X_text.toarray(), X_rep_pct))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```


Confusion Matrix for Logistic Regression:

Coefficients:


```
[[ 1.91723738e-02  5.28830096e-02 -1.40558302e-02 ... -5.77817776e-03
 -1.68158369e-01 -1.17107435e-04]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:




Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:



Out[17]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.647619 | 0.770186 | 0.756098 | 0.784810 | 0.230769 | 0.784810 | 0.507790 |
| 3 | Random Forest | 0.752381 | 0.857143 | 0.757282 | 0.987342 | 0.038462 | 0.987342 | 0.512902 |
| 4 | Perceptron | 0.533333 | 0.620155 | 0.800000 | 0.506329 | 0.615385 | 0.506329 | 0.560857 |
| 5 | KNN | 0.600000 | 0.727273 | 0.746667 | 0.708861 | 0.269231 | 0.708861 | 0.489046 |

Using 'Text' + male_jpc as the feature


```
In [18]: tfidf = TfidfVectorizer(min_df=5, max_df=0.8, ngram_range=(2, 2))
X_text = tfidf.fit_transform(df_rf['text'])
X_rep_pct = np.array(df_rf['rep_jpc']).reshape(-1, 1)
X_male_jpc = np.array(df_rf['male_jpc']).reshape(-1, 1)
X_dev_time = np.array(df_rf['develop_time']).reshape(-1, 1)
X_party = np.array(df_rf['party']).reshape(-1, 1)
y = df_rf['win_side']
# Horizontally stack the TF-IDF other features
X = np.hstack((X_text.toarray(), X_male_jpc))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```


Confusion Matrix for Logistic Regression:

Coefficients:


```
[[ 0.01916319  0.05289112 -0.01406044 ... -0.00579865 -0.16836936
 -0.0928783 ]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:




Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:



Out[18]:


| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.647619 | 0.770186 | 0.756098 | 0.784810 | 0.230769 | 0.784810 | 0.507790 |
| 3 | Random Forest | 0.685714 | 0.813559 | 0.734694 | 0.911392 | 0.000000 | 0.911392 | 0.455696 |
| 4 | Perceptron | 0.561905 | 0.676056 | 0.761905 | 0.607595 | 0.423077 | 0.607595 | 0.515336 |
| 5 | KNN | 0.609524 | 0.735484 | 0.750000 | 0.721519 | 0.269231 | 0.721519 | 0.495375 |

Using 'Text' + dev_time as the feature


```
In [29]: tfidf = TfidfVectorizer(min_df=5, max_df=0.8, ngram_range=(2, 2))
X_text = tfidf.fit_transform(df_rf['text'])
X_rep_pct = np.array(df_rf['rep_jpc']).reshape(-1, 1)
X_male_jpc = np.array(df_rf['male_jpc']).reshape(-1, 1)
X_dev_time = np.array(df_rf['develop_time']).reshape(-1, 1)
X_party = np.array(df_rf['party']).reshape(-1, 1)
y = df_rf['win_side']
# Horizontally stack the TF-IDF other features
X = np.hstack((X_text.toarray(), X_dev_time))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
```

Confusion Matrix for Logistic Regression:
Coefficients:
[[0.01901232 0.05250654 -0.01401245 ... -0.00578631 -0.16533421
-0.00045581]]




Confusion Matrix for Naive Bayes:




C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\svm_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(

Confusion Matrix for Linear SVC:




Confusion Matrix for Random Forest:




Confusion Matrix for Perceptron:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))



Confusion Matrix for KNN:



Out[29]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.257143 | 0.048780 | 0.666667 | 0.025316 | 0.961538 | 0.025316 | 0.493427 |
| 3 | Random Forest | 0.733333 | 0.840909 | 0.762887 | 0.936709 | 0.115385 | 0.936709 | 0.526047 |
| 4 | Perceptron | 0.247619 | 0.000000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 | 0.500000 |
| 5 | KNN | 0.676190 | 0.792683 | 0.764706 | 0.822785 | 0.230769 | 0.822785 | 0.526777 |


In [20]:

```
# USING IMBLEARN
from imblearn.over_sampling import SMOTE
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier(X_train, X_test, y_train, y_test)
ros = SMOTE(random_state=0)
X_train_resampled, y_train_resampled = ros.fit_resample(X_train, y_train)
Classifier(X_train_resampled, X_test, y_train_resampled, y_test)
```


Confusion Matrix for Logistic Regression:

Coefficients:

```
[[ 0.01901232  0.05250654 -0.01401245 ... -0.00578631 -0.16533421
 -0.00045581]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\svm_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(




Confusion Matrix for Random Forest:




Confusion Matrix for Perceptron:

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\metrics_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))




Confusion Matrix for KNN:




Confusion Matrix for Logistic Regression:

Coefficients:

```
[[ 0.01863534  0.06180091 -0.01704493 ... -0.00818723 -0.12273915
 -0.00084728]]
```




Confusion Matrix for Naive Bayes:




Confusion Matrix for Linear SVC:


C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\sklearn\svm_base.py:1244: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(




Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:



Out[20]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.647619 | 0.764331 | 0.769231 | 0.759494 | 0.307692 | 0.759494 | 0.533593 |
| 1 | Naive Bayes | 0.238095 | 0.000000 | 0.000000 | 0.000000 | 0.961538 | 0.000000 | 0.480769 |
| 2 | Linear SVC | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 3 | Random Forest | 0.657143 | 0.783133 | 0.747126 | 0.822785 | 0.153846 | 0.822785 | 0.488315 |
| 4 | Perceptron | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 5 | KNN | 0.466667 | 0.555556 | 0.744681 | 0.443038 | 0.538462 | 0.443038 | 0.490750 |

Check using different hyperparameter

```
In [57]: from sklearn.model_selection import GridSearchCV
vectorizer = TfidfVectorizer(min_df=5, max_df=0.7, ngram_range=(2, 2))
X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)

#Load hyperparam
param_grid = [
    {'penalty': ['l1'], 'solver': ['liblinear'], 'C': [0.1, 1, 10]},
    {'penalty': ['l2'], 'solver': ['lbfgs'], 'C': [0.1, 1, 10]},
    {'penalty': [None], 'solver': ['lbfgs']},
]

grid = GridSearchCV(LogisticRegression(max_iter=1000), param_grid, cv=10, scoring='f1',
                    return_train_score=True, error_score=np.nan)
grid.fit(X_train, y_train)
```

```

results = pd.DataFrame(grid.cv_results_)
results['param_penalty'] = results['param_penalty'].astype('str')
scores_matrix = results.pivot(index='param_penalty', columns='param_C', values='mean_test_score')

plt.figure(figsize=(8, 6))
sns.heatmap(scores_matrix, annot=True, fmt=".2f", linewidths=.5, cmap='viridis')
plt.title('Hyperparameter Gridsearch - f1 Score')
plt.show()


```

C:\Users\rezar\AppData\Local\Temp\ipykernel_10152\1042704739.py:19: FutureWarning: In a future version, the Index constructor will not infer numeric dtypes when passed object-dtype sequences (matching Series behavior)

```

scores_matrix = results.pivot(index='param_penalty', columns='param_C', values='mean_test_score')

```



Trying using pretrained model

Try using pretrained model from transformers.

https://huggingface.co/models?pipeline_tag=text-classification&sort=downloads

```

In [38]: import torch
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
# from transformers import BertTokenizer, BertForSequenceClassification, Trainer, TrainingArguments
from transformers import DistilBertTokenizer, DistilBertForSequenceClassification, Trainer, TrainingArguments

# Load your data
data = pd.read_csv('text_clean.csv')

# Swapped 0 and 1 in win_side
# data['win_side'] = data['win_side'].apply(lambda x: 0 if x == 1 else 1)

# Split your data into training and testing sets
train_texts, test_texts, train_labels, test_labels = train_test_split(data['text'], data['win_side'], test_size=0.2, stratify=data)

# Initialize the BERT tokenizer
# tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
tokenizer = DistilBertTokenizer.from_pretrained("distilbert-base-uncased")

# Tokenize the text data
train_encodings = tokenizer(train_texts.tolist(), truncation=True, padding=True)
test_encodings = tokenizer(test_texts.tolist(), truncation=True, padding=True)

# Create PyTorch Class from dataset
class SCOTUSDataset(torch.utils.data.Dataset):
    def __init__(self, encodings, labels):

```

```

        self.encodings = encodings
        self.labels = labels

    def __getitem__(self, idx):
        item = {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
        item['labels'] = torch.tensor(self.labels[idx]).long()
        return item

    def __len__(self):
        return len(self.labels)

train_dataset = SCOTUSDataset(train_encodings, train_labels.tolist())
test_dataset = SCOTUSDataset(test_encodings, test_labels.tolist())

# Initialize the BERT model for sequence classification
# model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=2)
# https://huggingface.co/distilbert-base-uncased-finetuned-sst-2-english
model = DistilBertForSequenceClassification.from_pretrained("distilbert-base-uncased")

# Set up training arguments
training_args = TrainingArguments(
    output_dir='./results',
    num_train_epochs=10,
    per_device_train_batch_size=2, #changed to 2 because the GPU
    per_device_eval_batch_size=4,
    warmup_steps=500,
    weight_decay=0.01,
    logging_dir='./logs',
    logging_steps=100,
    evaluation_strategy="steps",
    save_strategy="steps",
    save_steps=1000,
    load_best_model_at_end=True,
)

# Create the Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=test_dataset
)

# Train the model
trainer.train()

# Evaluate the model
predictions = trainer.predict(test_dataset)
predicted_labels = predictions.label_ids
y_pred = (predictions.predictions.argmax(-1)).tolist()

# Calculate performance metrics
accuracy = accuracy_score(test_labels, y_pred)
precision = precision_score(test_labels, y_pred)
recall = recall_score(test_labels, y_pred)
f1 = f1_score(test_labels, y_pred)

print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")

```

Some weights of the model checkpoint at distilbert-base-uncased were not used when initializing DistilBertForSequenceClassification: ['vocab_projector.weight', 'vocab_projector.bias', 'vocab_layer_norm.bias', 'vocab_transform.bias', 'vocab_transform.weight', 'vocab_layer_norm.weight']

- This IS expected if you are initializing DistilBertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing DistilBertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['pre_classifier.bias', 'classifier.weight', 'classifier.bias', 'pre_classifier.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

C:\Users\rezar\anaconda3\envs\torch_cuda\lib\site-packages\transformers\optimization.py:391: FutureWarning: This implementation of AdamW is deprecated and will be removed in a future version. Use the PyTorch implementation torch.optim.AdamW instead, or set `no_deprecation_warning=True` to disable this warning
warnings.warn()

| Step | Training Loss | Validation Loss |
|------|---------------|-----------------|
| 100 | 0.694300 | 0.668718 |
| 200 | 0.664300 | 0.656595 |
| 300 | 0.687600 | 0.654255 |
| 400 | 0.679200 | 0.660695 |
| 500 | 0.673700 | 0.654594 |
| 600 | 0.694700 | 0.651798 |
| 700 | 0.651100 | 0.674763 |
| 800 | 0.680400 | 0.689492 |
| 900 | 0.686100 | 0.665781 |
| 1000 | 0.670000 | 0.656918 |
| 1100 | 0.672700 | 0.657205 |
| 1200 | 0.656400 | 0.654485 |
| 1300 | 0.676800 | 0.665568 |
| 1400 | 0.671300 | 0.656507 |
| 1500 | 0.634500 | 0.665914 |
| 1600 | 0.684400 | 0.656613 |
| 1700 | 0.665900 | 0.655252 |
| 1800 | 0.660500 | 0.654917 |
| 1900 | 0.671700 | 0.654646 |
| 2000 | 0.632400 | 0.656990 |

Accuracy: 0.6381
Precision: 0.6381
Recall: 1.0000
F1-Score: 0.7791

```
In [59]: def Classifier2(X_train, X_test, y_train, y_test):

    # Train and evaluate the classifiers
    classifiers = {
        "Logistic Regression": LogisticRegression(max_iter=1000, penalty='l2', C=1, solver='lbfgs'),
        "Naive Bayes": MultinomialNB(),
        "Linear SVC": LinearSVC(),
        "Random Forest": RandomForestClassifier(),
        "Perceptron": Perceptron(),
        "KNN": KNeighborsClassifier(n_neighbors=11)
    }
    results = []

    for classifier_name, classifier in classifiers.items():

        # Train the classifier
        classifier.fit(X_train, y_train)

        # Make predictions on the test set
        y_pred = classifier.predict(X_test)

        # Add the scores to the results dictionary
        results.append({
            'classifier': classifier_name,
            'accuracy': accuracy_score(y_test, y_pred),
            'f1': f1_score(y_test, y_pred),
            'precision': precision_score(y_test, y_pred),
            'recall': recall_score(y_test, y_pred),
            'True Negative Rate': confusion_matrix(y_test, y_pred)[0][0]/(confusion_matrix(y_test, y_pred)[0][0]+confusion_matrix(
            'True Positive Rate': confusion_matrix(y_test, y_pred)[1][1]/(confusion_matrix(y_test, y_pred)[1][1]+confusion_matrix(
            'auc': roc_auc_score(y_test, y_pred)
        })

        # Make a confusion matrix
        print(f"Confusion Matrix for {classifier_name}:")

        if classifier_name == "Logistic Regression":
```

```

print("Coefficients: \n", classifier.coef_)

cm = confusion_matrix(y_test, y_pred, labels=y_test.unique())
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=y_test.unique())
fig, ax = plt.subplots(figsize=(3, 3))
ax.set_title(f"CM {classifier_name}:")
disp.plot(ax=ax)
plt.show()
return pd.DataFrame(results)

```

```

In [60]: X, y = Vectorize(vectorizer, df_rf['text'], df_rf['win_side'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=55)
Classifier2(X_train, X_test, y_train, y_test)

```


Confusion Matrix for Logistic Regression:

Coefficients:


```

[[ 0.01914151  0.05325384 -0.01408864 ... -0.01701203 -0.00579603
 -0.16882371]]


```




Confusion Matrix for Naïve Bayes:




Confusion Matrix for Linear SVC:




Confusion Matrix for Random Forest:



Confusion Matrix for Perceptron:



Confusion Matrix for KNN:



Out[60]:

| | classifier | accuracy | f1 | precision | recall | True Negative Rate | True Positive Rate | auc |
|---|---------------------|----------|----------|-----------|----------|--------------------|--------------------|----------|
| 0 | Logistic Regression | 0.761905 | 0.863388 | 0.759615 | 1.000000 | 0.038462 | 1.000000 | 0.519231 |
| 1 | Naive Bayes | 0.752381 | 0.858696 | 0.752381 | 1.000000 | 0.000000 | 1.000000 | 0.500000 |
| 2 | Linear SVC | 0.647619 | 0.770186 | 0.756098 | 0.784810 | 0.230769 | 0.784810 | 0.507790 |
| 3 | Random Forest | 0.733333 | 0.846154 | 0.747573 | 0.974684 | 0.000000 | 0.974684 | 0.487342 |
| 4 | Perceptron | 0.600000 | 0.716216 | 0.768116 | 0.670886 | 0.384615 | 0.670886 | 0.527751 |
| 5 | KNN | 0.685714 | 0.795031 | 0.780488 | 0.810127 | 0.307692 | 0.810127 | 0.558909 |

In []: